

DynamoDB

Design Patterns and Best Practices

Rick Houlihan, Principal Solutions Architect


1/20/2016

What to expect from the session



- Brief history of data processing
- DynamoDB Internals
 - Tables, API, data types, indexes
 - Scaling and data modeling
- Design patterns and best practices
- Event driven applications and DDB Streams

Timeline of Database Technology


Data Pressure



Ledgers




Unit Records





Data Drums




File Systems




RDBMS



NoSQL




Data Volume Since 2010



- 90% of stored data generated in last 2 years
- 1 Terabyte of data in 2010 equals 6.5 Petabytes today
- Linear correlation between data pressure and technical innovation
- No reason these trends will not continue over time

Technology Adoption and the Hype Cycle



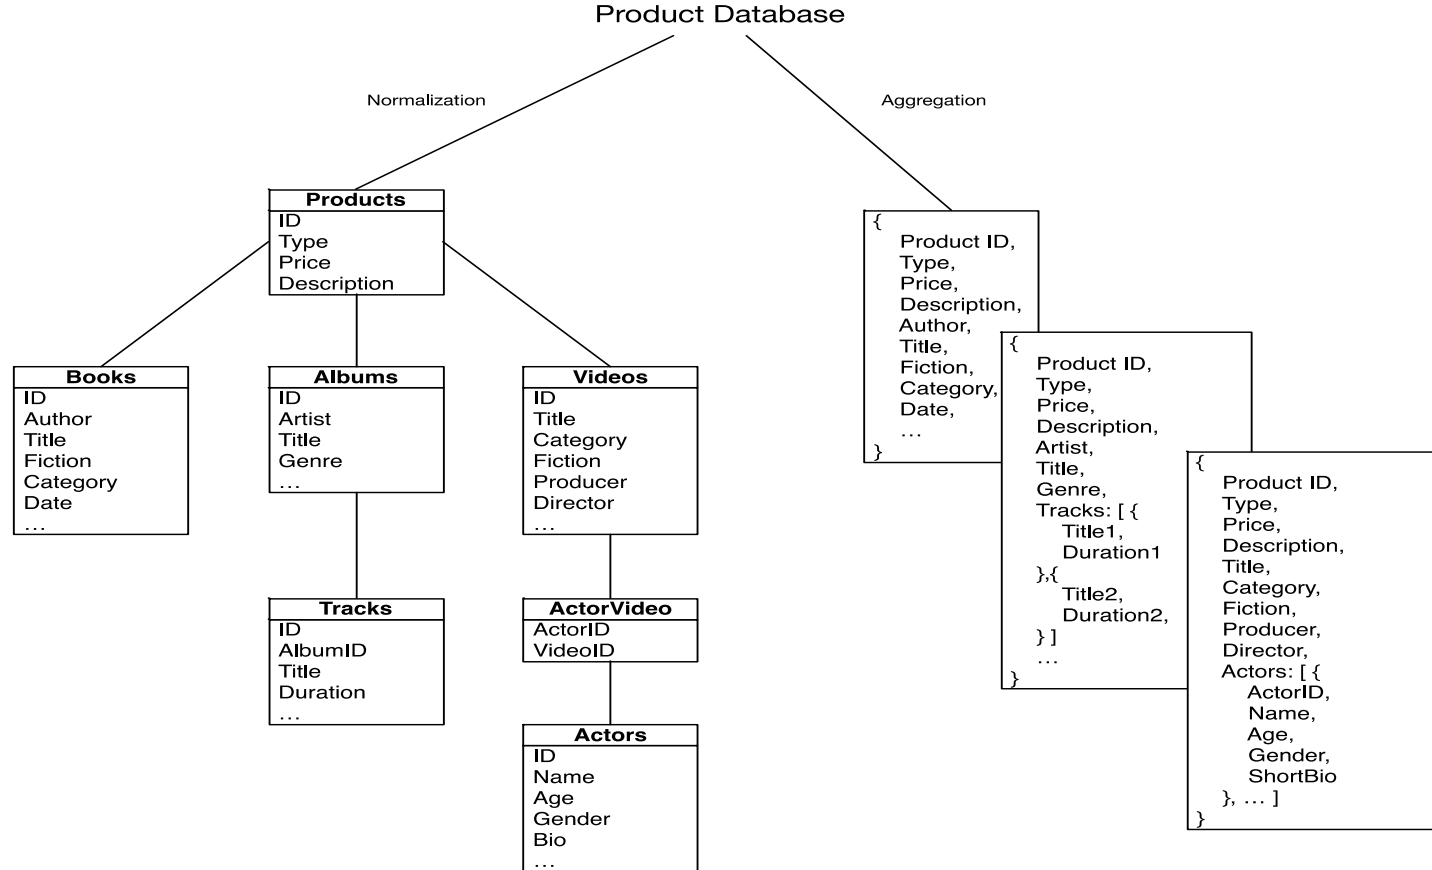
Why NoSQL?

SQL


NoSQL

Optimized for storage	Optimized for compute
Normalized/relational	Denormalized/hierarchical
Ad hoc queries	Instantiated views
Scale vertically	Scale horizontally
Good for OLAP	Built for OLTP at scale

SQL vs. NoSQL Access Pattern



Table



Partition
Key

Sort
Key

Mandatory
Key-value access pattern
Determines data distribution

Optional
Model 1:N relationships
Enables rich query capabilities

All items for key
==, <, >, >=, <=
"begins with"
"between"
"contains"
"in"
sorted results
counts
top/bottom N values

Partition Keys

Partition Key uniquely identifies an item

Partition Key is used for building an unordered hash index

Allows table to be partitioned for scale

Id = 1
Name = Jim


Hash (1) = 7B

Id = 2
Name = Andy
Dept = Eng

Hash (2) = 48

Id = 3
Name = Kim
Dept = Ops

Hash (3) = CD



Partition:Sort Key

Partition:Sort Key uses two attributes together to uniquely identify an Item

Within unordered hash index, data is arranged by the sort key

No limit on the number of items (∞) per partition key

- Except if you have local secondary indexes

Partition 1

Partition 2

Partition 3

00:0


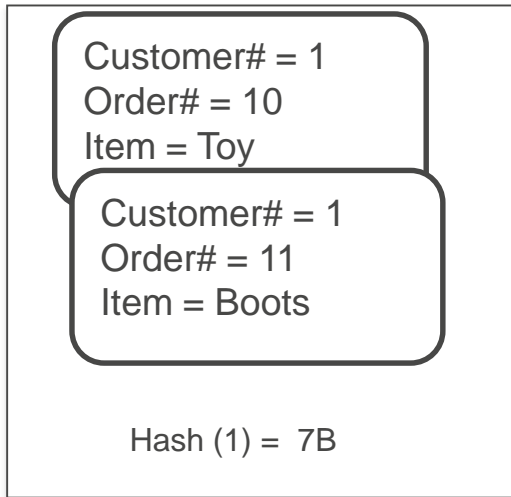
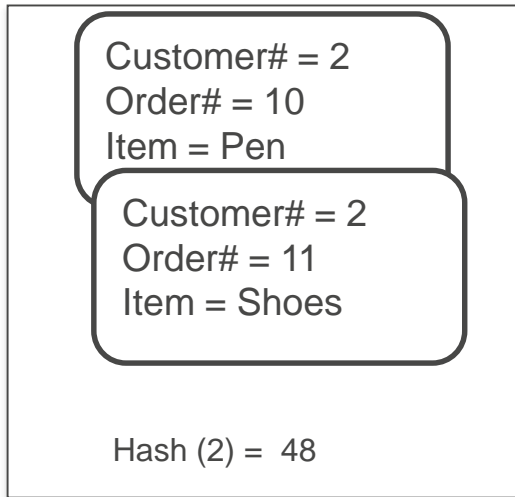
54: ∞

55

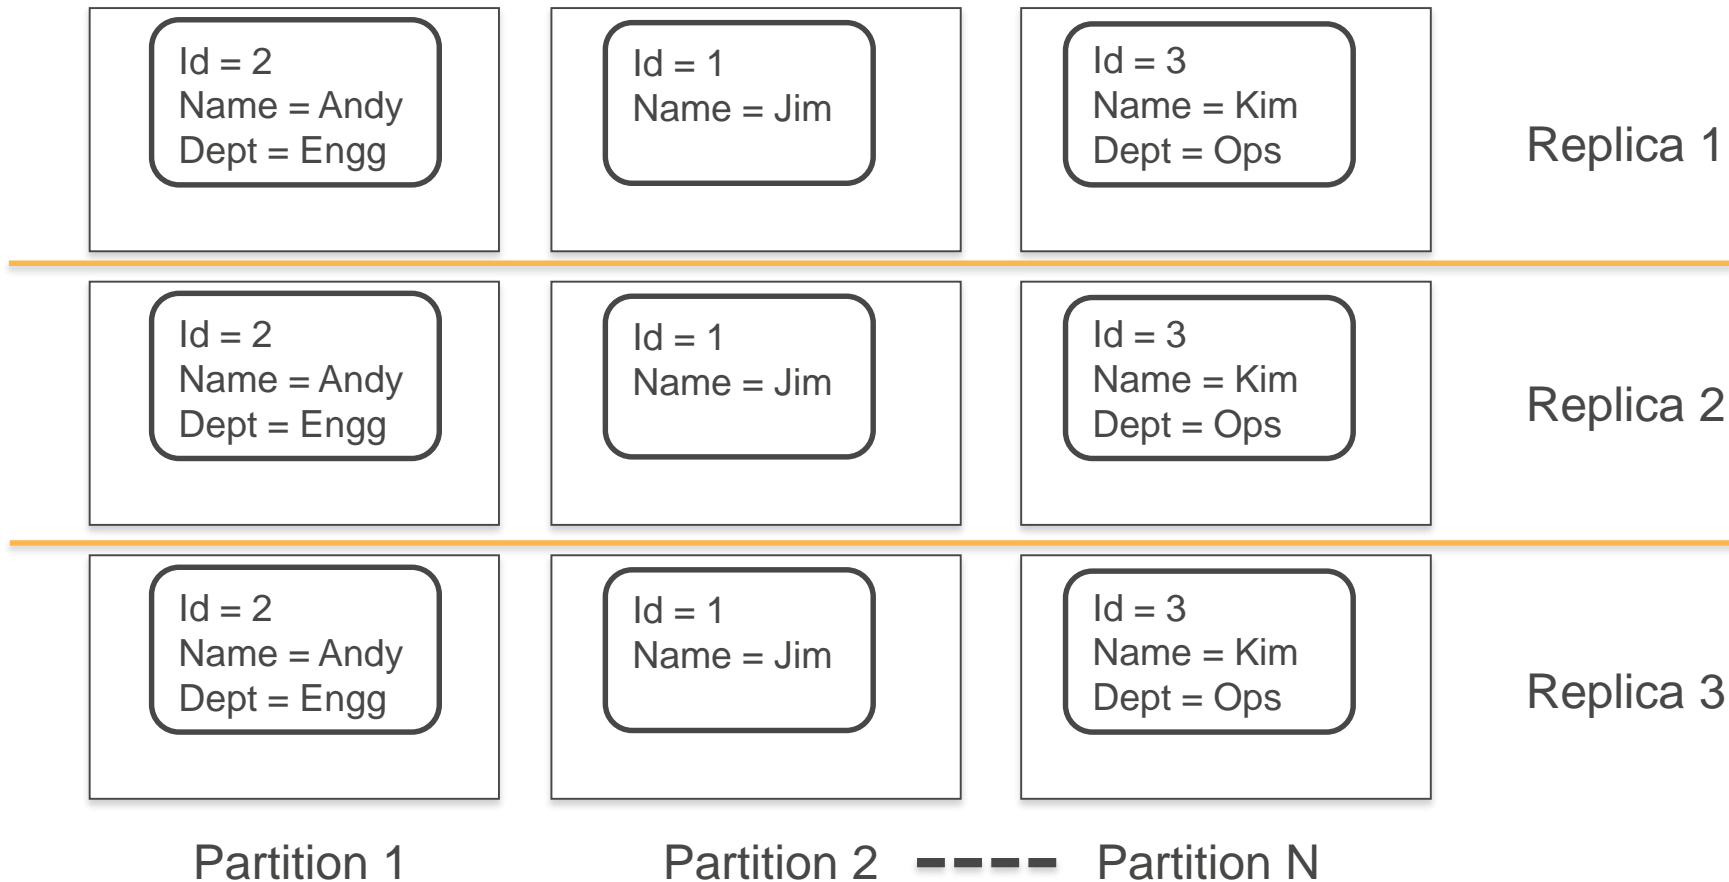
A9: ∞

AA

FF: ∞



Partitions are three-way replicated




Indexes

Local secondary index (LSI)

Alternate sort key attribute

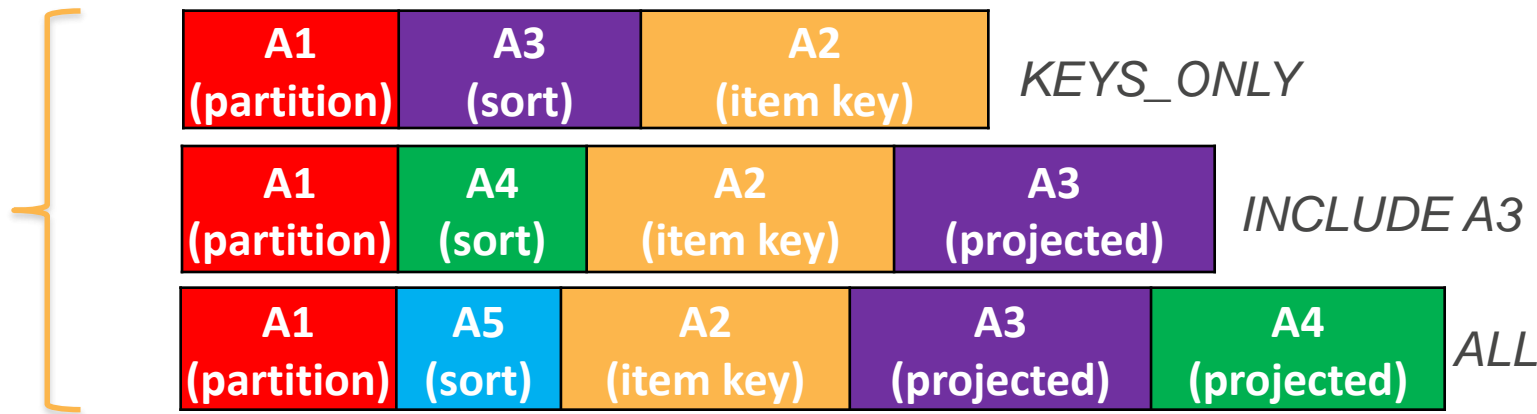
Index is local to a partition key

Table



10 GB max per partition
key, i.e. LSIs limit the # of
range keys!

LSIs

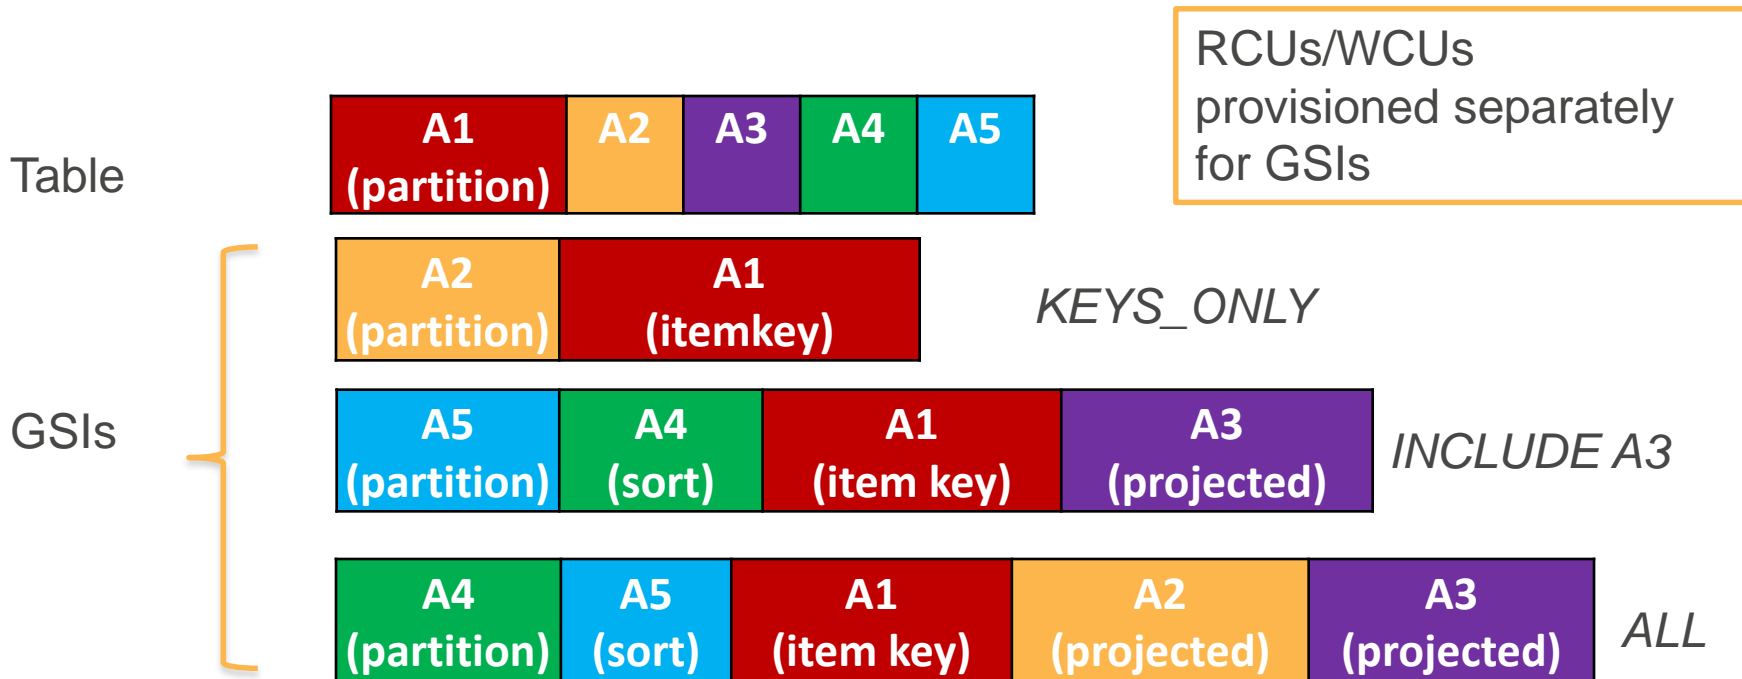


Global secondary index (GSI)

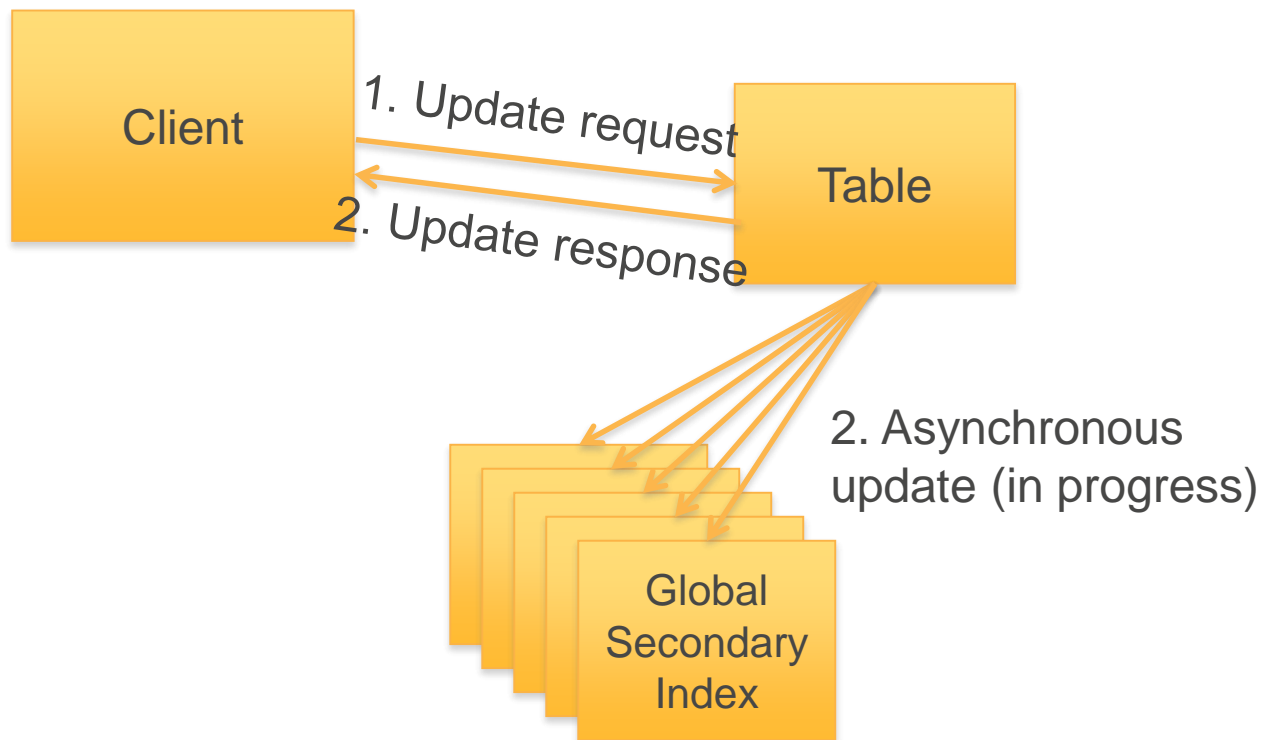
Alternate partition and/or sort key

Index is across all partition keys

Online indexing



How do GSI updates work?



If GSIs don't have enough write capacity, table writes will be throttled!

LSI or GSI?

LSI can be modeled as a GSI

If data size in an item collection > 10 GB, use GSI

If eventual consistency is okay for your scenario, use GSI!

Scaling

Scaling

Throughput

- Provision any amount of throughput to a table

Size

- Add any number of items to a table
 - Max item size is 400 KB
 - LSIs limit the number of range keys due to 10 GB limit

Scaling is achieved through partitioning

Throughput

Provisioned at the table level

- Write capacity units (WCUs) are measured in 1 KB per second
- Read capacity units (RCUs) are measured in 4 KB per second
 - RCUs measure strictly consistent reads
 - Eventually consistent reads cost 1/2 of consistent reads

Read and write throughput limits are independent



RCU



WCU

Partitioning math

Number of Partitions	
By Capacity	$(\text{Total RCU} / 3000) + (\text{Total WCU} / 1000)$
By Size	Total Size / 10 GB
Total Partitions	$\text{CEILING}(\text{MAX}(\text{Capacity}, \text{Size}))$

In the future, these details might change...

Partitioning example

Table size = 8 GB, RCUs = 5000, WCUs = 500

Number of Partitions	
By Capacity	$(5000 / 3000) + (500 / 1000) = 2.17$
By Size	$8 / 10 = 0.8$
Total Partitions	$\text{CEILING}(\text{MAX}(2.17, 0.8)) = 3$

RCUs and WCUs are uniformly spread across partitions

RCUs per partition = $5000/3 = 1666.67$
WCUs per partition = $500/3 = 166.67$
Data/partition = $10/3 = 3.33$ GB

What causes throttling?

If **sustained** throughput goes beyond provisioned throughput per partition

Non-uniform workloads

- Hot keys/hot partitions
- Very large bursts

Mixing hot data with cold data


- Use a table per time period



From the example before:

- Table created with 5000 RCUs, 500 WCUs
- RCUs per partition = 1666.67
- WCUs per partition = 166.67
- If sustained throughput > (1666 RCUs or 166 WCUs) per key or partition, DynamoDB may throttle requests
 - Solution: Increase provisioned throughput

What bad NoSQL looks like...



Getting the most out of DynamoDB throughput


“To get the most out of DynamoDB throughput, create tables where the hash key element has a large number of distinct values, and values are requested fairly uniformly, as randomly as possible.”

—*DynamoDB Developer Guide*

Space: access is evenly spread over the key-space

Time: requests arrive evenly spaced in time

Much better picture...



Data modeling

1:1 relationships or key-values

Use a table or GSI with an alternate partition key

Use GetItem or BatchGetItem API

Example: Given an SSN or license number, get attributes

Users Table	
Partiton key	Attributes
SSN = 123-45-6789	Email = johndoe@nowhere.com, License = TDL25478134
SSN = 987-65-4321	Email = maryfowler@somewhere.com, License = TDL78309234

Users-Email-GSI	
Partition key	Attributes
License = TDL78309234	Email = maryfowler@somewhere.com, SSN = 987-65-4321
License = TDL25478134	Email = johndoe@nowhere.com, SSN = 123-45-6789

1:N relationships or parent-children

Use a table or GSI with partition and sort key

Use Query API

Example:

- Given a device, find all readings between epoch X, Y

Device-measurements		
Partition Key	Sort key	Attributes
DeviceId = 1	epoch = 5513A97C	Temperature = 30, pressure = 90
DeviceId = 1	epoch = 5513A9DB	Temperature = 30, pressure = 90

N:M relationships

Use a table and GSI with partition and sort key elements switched



Use Query API

Example: Given a user, find all games. Or given a game, find all users.

User-Games-Table	
Partition Key	Sort key
UserId = bob	GameId = Game1
UserId = fred	GameId = Game2
UserId = bob	GameId = Game3

Game-Users-GSI	
Partition Key	Sort key
GameId = Game1	UserId = bob
GameId = Game2	UserId = fred
GameId = Game3	UserId = bob

Hierarchical Data



Tiered relational data structures

Hierarchical Data Structures as Items...

Use composite sort key to define a Hierarchy

Highly selective result sets with sort queries

Index anything, scales to any size

	Primary Key		Attributes								
	ProductID	type									
Items	1	bookID	title	author	genre	publisher	datePublished	ISBN			
			Ringworld	Larry Niven	Science Fiction	Ballantine	Oct-70	0-345-02046-4			
	2	albumID	title	artist	genre	label	studio	releasd	producer		
			Dark Side of the Moon	Pink Floyd	Progressive Rock	Harvest	Abbey Road	3/1/73	Pink Floyd		
	2	albumID:trackID	title	length	music	vocals					
			Speak to Me	1:30	Mason	Instrumental					
	2	albumID:trackID	title	length	music	vocals					
			Breathe	2:43	Waters, Gilmour, Wright	Gilmour					
	2	albumID:trackID	title	length	music	vocals					
			On the Run	3:30	Gilmour, Waters	Instrumental					
	3	movieID	title	genre	writer	producer					
			Idiocracy	Scifi Comedy	Mike Judge	20th Century Fox					
	3	movieID:actorID	name	character	image						
			Luke Wilson	Joe Bowers	img2.jpg						
	3	movieID:actorID	name	character	image						
			Maya Rudolph	Rita	img3.jpg						
	3	movieID:actorID	name	character	image						
			Dax Shepard	Frito Pendejo	img1.jpg						

... or as Documents (JSON)

JSON data types (M, L, BOOL, NULL)

Document SDKs Available



Indexing only via Streams/Lambda

400KB max item size (limits hierarchical data structure)

	Primary Key		Attributes						
	ProductID								
Items	1	id	title	author	genre	publisher	datePublished	ISBN	
		bookID	Ringworld	Larry Niven	Science Fiction	Ballantine	Oct-70	0-345-02046-4	
	2	id	title	artist	genre	Attributes			
		albumID	Dark Side of the Moon	Pink Floyd	Progressive Rock	{ label: "Harvest", studio: "Abbey Road", published: "3/1/73", producer: "Pink Floyd", tracks: [{title: "Speak to Me", length: "1:30", music: "Mason", vocals: "Instrumental"}], {title: "Breathe", length: "2:43", music: "Waters, Gilmour, Wright", vocals: "Gilmour"}], {title: "On the Run", length: "3:30", music: "Gilmour, Waters", vocals: "Instrumental"}]}			
	3	id	title	genre	writer	Attributes			
		movieID	Idiocracy	Scifi Comedy	Mike Judge	{ producer: "20th Century Fox", actors: [{ name: "Luke Wilson", dob: "9/21/71", character: "Joe Bowers", image: "img2.jpg" }, { name: "Maya Rudolph", dob: "7/27/72", character: "Rita", image: "img1.jpg" }, { name: "Dax Shepard", dob: "1/2/75", character: "Frito Pendejo", image: "img3.jpg" }]}			

Scenarios and best practices

Event logging



Storing time series data

Time series tables

Current table

Events_table_2015_April				
<u>Event_id</u> (Partition)	<u>Timestamp</u> (Sort)	Attribute1	Attribute N

RCUs = 10000
WCUs = 10000

Older tables

Events_table_2015_March				
<u>Event_id</u> (Partition)	<u>Timestamp</u> (Sort)	Attribute1	Attribute N

RCUs = 1000
WCUs = 1

Events_table_2015_Febuary				
<u>Event_id</u> (Partition)	<u>Timestamp</u> (Sort)	Attribute1	Attribute N

RCUs = 100
WCUs = 1

Events_table_2015_January				
<u>Event_id</u> (Partition)	<u>Timestamp</u> (Sort)	Attribute1	Attribute N

RCUs = 10
WCUs = 1

Hot data

Cold data

Don't mix hot and cold data; archive cold data to Amazon S3



Use a table per time period

Pre-create daily, weekly, monthly tables




Provision required throughput for current table

Writes go to the current table

Turn off (or reduce) throughput for older tables

Important when: Dealing with time series data

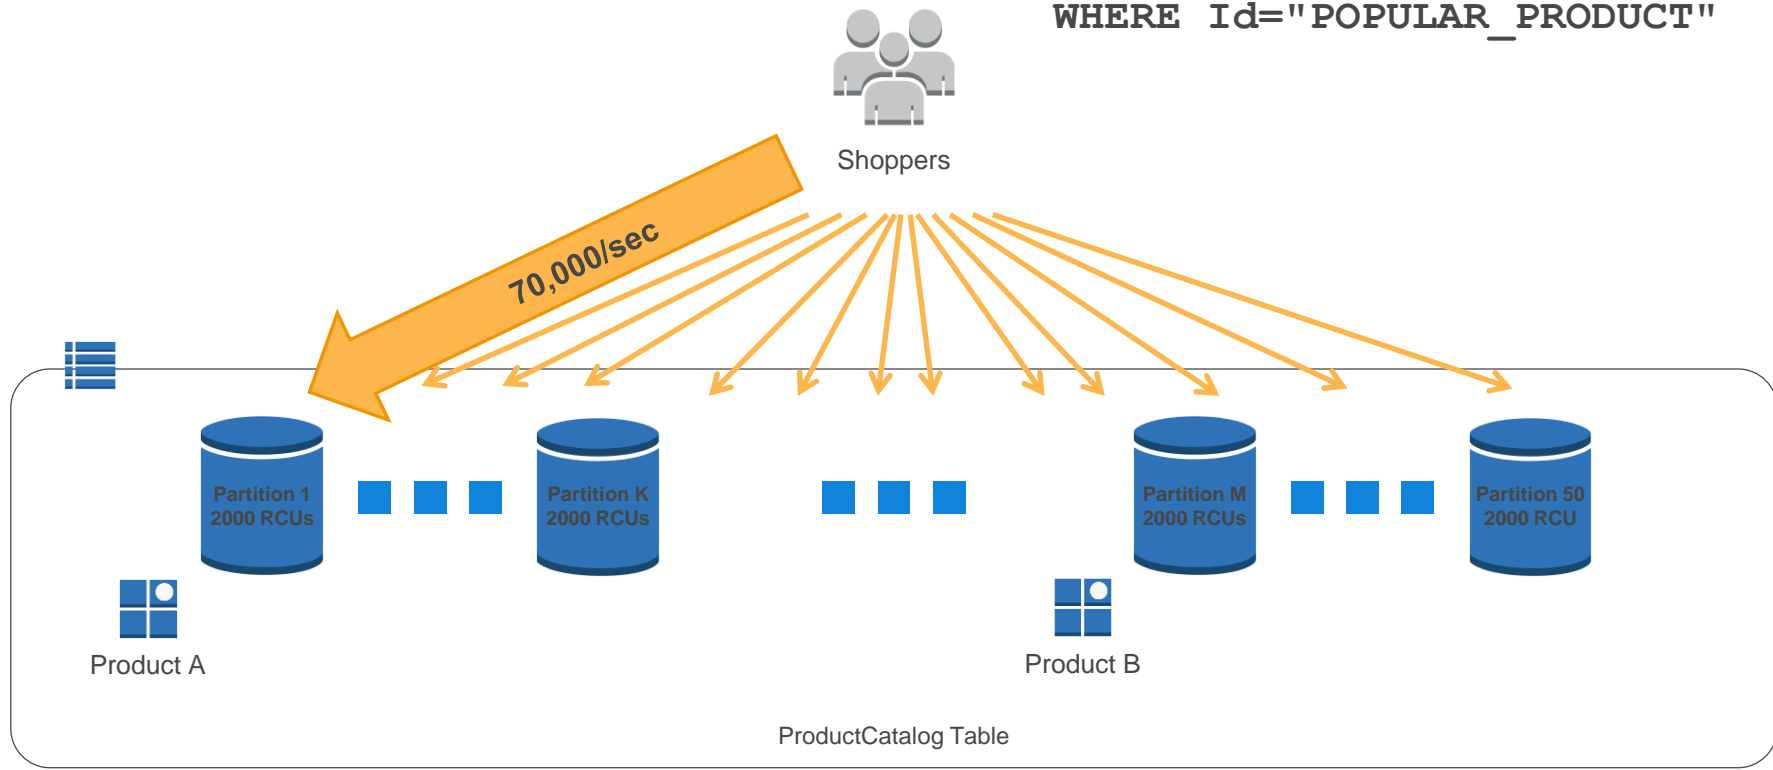
Product catalog




Popular items (read)

Scaling bottlenecks

```
SELECT Id, Description, ...  
FROM ProductCatalog  
WHERE Id="POPULAR_PRODUCT"
```




Request Distribution Per Partition Key




Cache popular items



```
SELECT Id, Description, ...  
FROM ProductCatalog  
WHERE Id="POPULAR_PRODUCT"
```



Request Distribution Per Partition Key



Messaging app



Large items

Filters vs. indexes

M:N Modeling—inbox and outbox



David



Messages App



Messages Table

Inbox

```
SELECT *  
FROM Messages  
WHERE Recipient='David'  
LIMIT 50  
ORDER BY Date DESC
```

Outbox

```
SELECT *  
FROM Messages  
WHERE Sender = 'David'  
LIMIT 50  
ORDER BY Date DESC
```

Large and small attributes mixed



David

Inbox

```
SELECT *  
FROM Messages  
WHERE Recipient='David'  
LIMIT 50  
ORDER BY Date DESC
```

Partition key

Sort key



Messages Table

Recipient	Date	Sender	Message
David	2014-10-02	Bob	...
... 48 more messages for David ...			
David	2014-10-03	Alice	...
Alice	2014-09-28	Bob	...
Alice	2014-10-01	Carol	...

50 items × 256 KB each

Large message bodies
Attachments

(Many more messages)

Computing inbox query cost

$$50 * 256\text{KB} * (1 \text{ RCU} / 4\text{KB}) * (1 / 2) = 1600 \text{ RCU}$$

Items evaluated by query

Average item size

Conversion ratio

Eventually consistent reads

Separate the bulk data

Uniformly distributes large item reads



David

(50 sequential items at 128 bytes)

1. Query Inbox-GSI: 1 RCU
2. BatchGetItem Messages: 1600 RCU

(50 separate items at 256 KB)

 Inbox-GSI

<u>Recipient</u>	<u>Date</u>	Sender	Subject	MsgId
David	2014-10-02	Bob	Hi!...	afed
David	2014-10-03	Alice	RE: The...	3kf8
Alice	2014-09-28	Bob	FW: Ok...	9d2b
Alice	2014-10-01	Carol	Hi!...	ct7r

 Messages Table

<u>MsgId</u>	Body
9d2b	...
3kf8	...
ct7r	...
afed	...

Inbox GSI

Define which attributes to copy into the index

Details	Indexes	Monitoring	Alarm Setup
Global Secondary Indexes			
Index Name	Hash Key	Range Key	Projected Attributes
Inbox	Recipient (String)	Date (String)	MsgId, Recipient, Date, Subject, Sender

Outbox GSI

Details

Indexes

Monitoring


Alarm Setup

Global Secondary Indexes

Index Name	Hash Key	Range Key	Projected Attributes
Outbox	Sender (String)	Date (String)	MsgId, Recipient, Date, Subject, Sender

```
SELECT *  
FROM Messages  
WHERE Sender = 'David'  
LIMIT 50  
ORDER BY Date DESC
```


Messaging app






Distribute large items



Reduce one-to-many item sizes

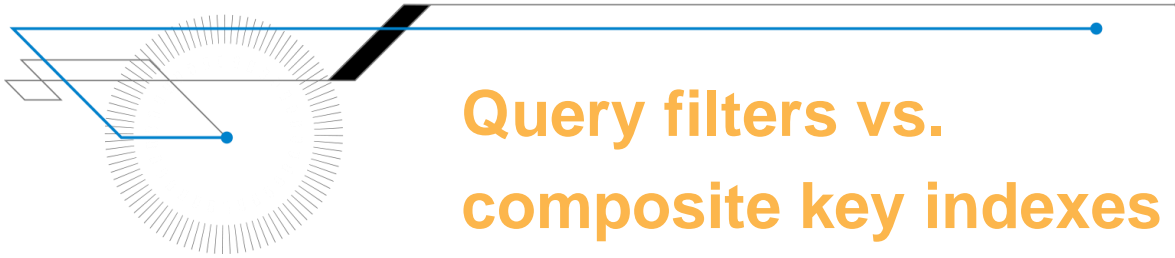

Configure secondary index projections

Use GSIs to model M:N relationship
between sender and recipient



Important when: Querying many large items at once

Multiplayer online gaming



Query filters vs.
composite key indexes

Hierarchical Data Structures

Partition key

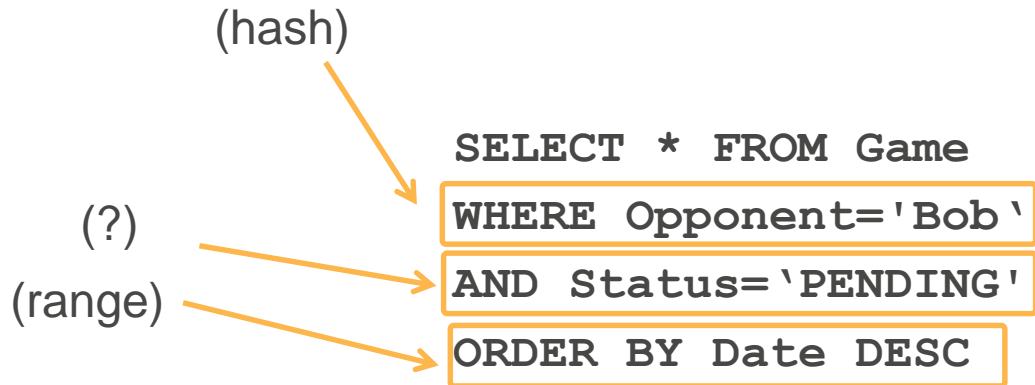
Games Table

<u>Gameld</u>	Date	Host	Opponent	Status
d9bl3	2014-10-02	David	Alice	DONE
72f49	2014-09-30	Alice	Bob	PENDING
o2pnb	2014-10-08	Bob	Carol	IN_PROGRESS
b932s	2014-10-03	Carol	Bob	PENDING
ef9ca	2014-10-03	David	Bob	IN_PROGRESS

Query for incoming game requests

DynamoDB indexes provide partition and sort

What about queries for two equalities and a sort?



Approach 1: Query filter



Bob

Partition key

Sort key



Secondary Index

<u>Opponent</u>	<u>Date</u>	<u>GameId</u>	Status	Host
Alice	2014-10-02	d9bl3	DONE	David
Carol	2014-10-08	o2pnb	IN_PROGRESS	Bob
Bob	2014-09-30	72f49	PENDING	Alice
Bob	2014-10-03	b932s	PENDING	Carol
Bob	2014-10-03	ef9ca	IN_PROGRESS	David

Approach 1: Query filter

```
SELECT * FROM Game
WHERE Opponent='Bob'
ORDER BY Date DESC
FILTER ON Status='PENDING'
```



Bob




Secondary Index

<u>Opponent</u>	<u>Date</u>	<u>GameId</u>	Status	Host
Alice	2014-10-02	d9bl3	DONE	David
Carol	2014-10-08	o2pnb	IN_PROGRESS	Bob
Bob	2014-09-30	72f49	PENDING	Alice
Bob	2014-10-03	b932s	PENDING	Carol
Bob	2014-10-03	ef9ca	IN_PROGRESS	David


← (filtered out)

Needle in a haystack





Use query filter



Send back less data “on the wire”

Simplify application code

Simple SQL-like expressions

- AND, OR, NOT, ()

Important when: Your index isn't entirely selective

Approach 2: Composite key

Status		Date		StatusDate
DONE		2014-10-02		DONE_2014-10-02
IN_PROGRESS		2014-10-08		IN_PROGRESS_2014-10-08
IN_PROGRESS	+	2014-10-03	=	IN_PROGRESS_2014-10-03
PENDING		2014-10-03		PENDING_2014-09-30
PENDING		2014-09-30		PENDING_2014-10-03

Approach 2: Composite key

Partition key

Sort key



Secondary Index

<u>Opponent</u>	<u>StatusDate</u>	<u>GameId</u>	Host
Alice	DONE_2014-10-02	d9bl3	David
Carol	IN_PROGRESS_2014-10-08	o2pnb	Bob
Bob	IN_PROGRESS_2014-10-03	ef9ca	David
Bob	PENDING_2014-09-30	72f49	Alice
Bob	PENDING_2014-10-03	b932s	Carol

Approach 2: Composite key

```
SELECT * FROM Game
WHERE Opponent='Bob'
AND StatusDate BEGINS_WITH 'PENDING'
```




Bob

 Secondary Index

<u>Opponent</u>	<u>StatusDate</u>	<u>GameId</u>	Host
Alice	DONE_2014-10-02	d9bl3	David
Carol	IN_PROGRESS_2014-10-08	o2pnb	Bob
Bob	IN_PROGRESS_2014-10-03	ef9ca	David
Bob	PENDING_2014-09-30	72f49	Alice
Bob	PENDING_2014-10-03	b932s	Carol

Needle in a *sorted* haystack



Sparse indexes

Game-scores-table



Id (Partition)	User	Game	Score	Date	Award
1	Bob	G1	1300	2012-12-23	
2	Bob	G1	1450	2012-12-23	
3	Jay	G1	1600	2012-12-24	
4	Mary	G1	2000	2012-10-24	Champ
5	Ryan	G2	123	2012-03-10	
6	Jones	G2	345	2012-03-20	

Scan sparse GSIs

Award-GSI

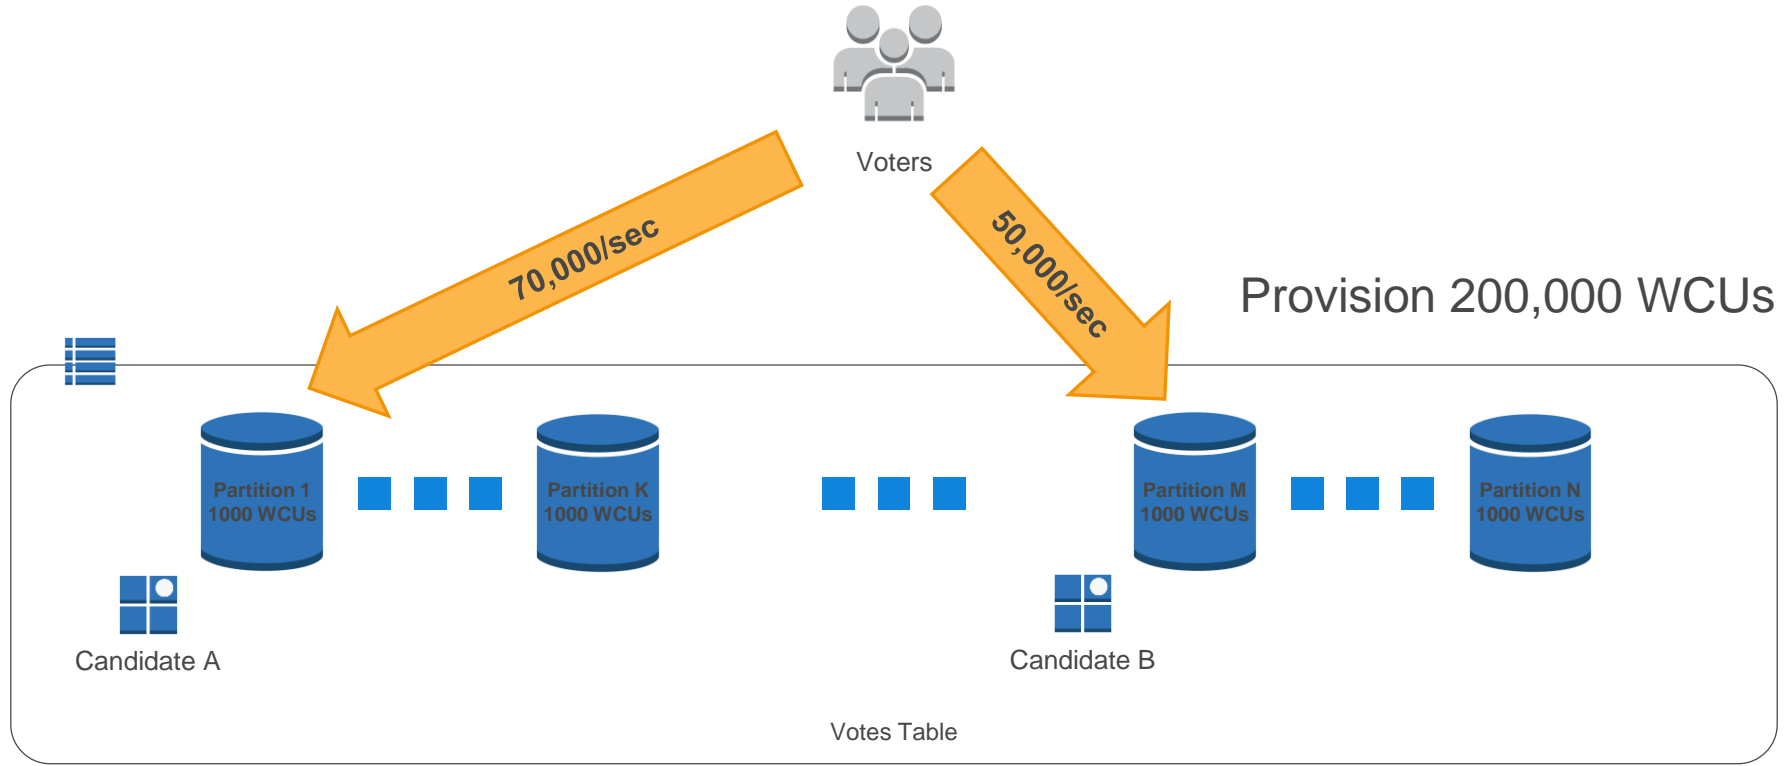
Award (Partition)	Id	User	Score
Champ	4	Mary	2000

Real-Time voting



Write-heavy items

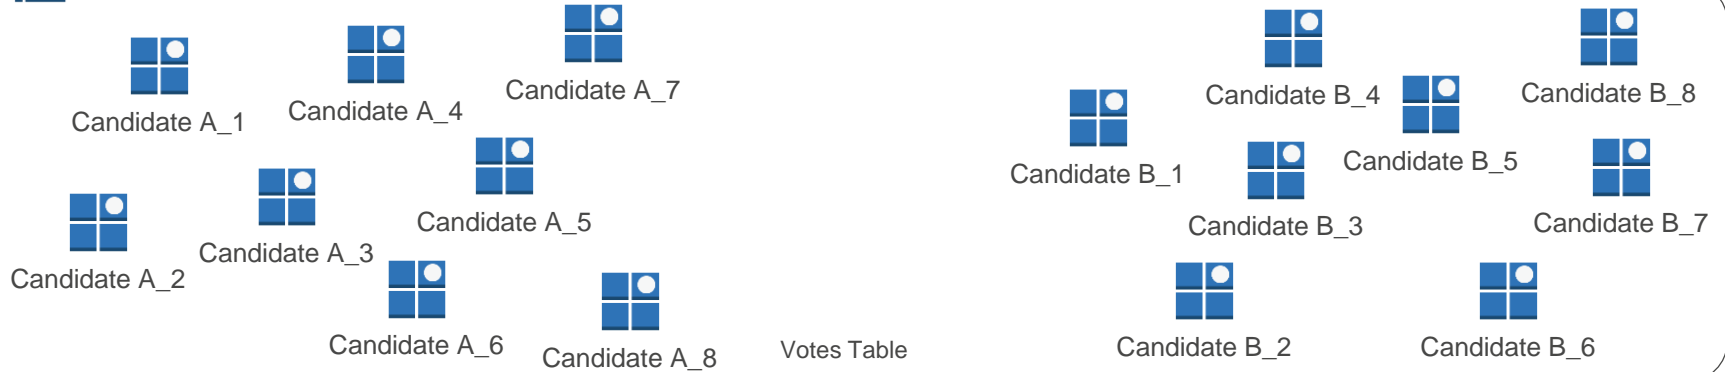
Scaling bottlenecks



Write sharding



Voter




Write sharding

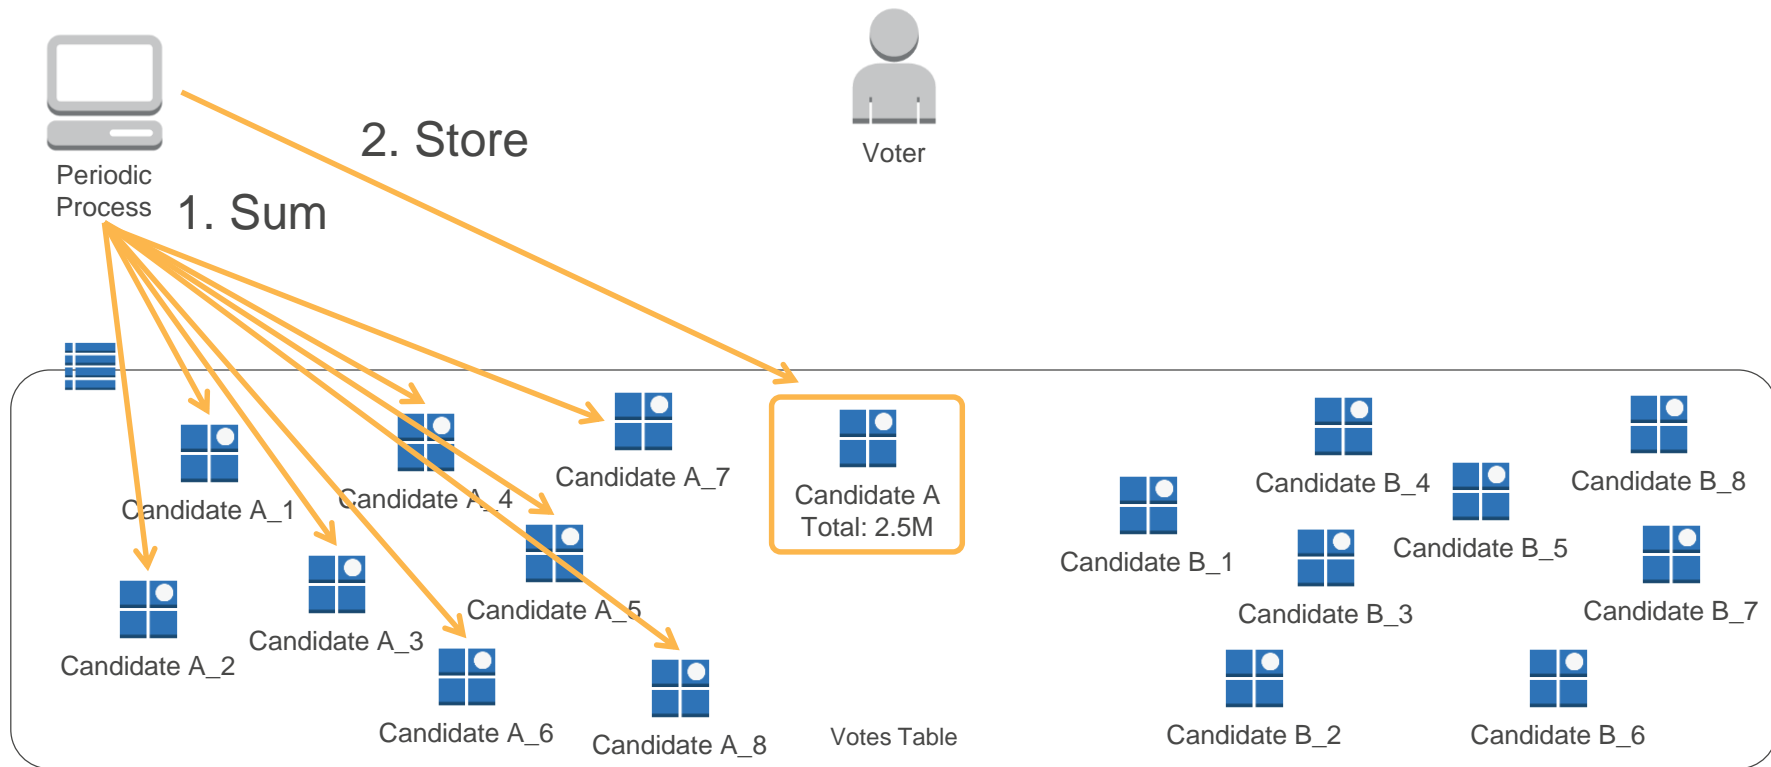


Voter

UpdateItem: "**CandidateA_**" + $\text{rand}(0, 10)$
ADD 1 to Votes




Shard aggregation






Shard write-heavy partition keys



Trade off read cost for write scalability

Consider throughput per partition key



Important when: Your write workload is not horizontally scalable



Replace filter with indexes



Concatenate attributes to form useful
secondary index keys




Take advantage of sparse indexes

Important when: You want to optimize a query as much
as possible

DynamoDB Streams

DynamoDB Streams



Stream of updates to a table

Asynchronous

Exactly once

Strictly ordered

- Per item

Highly durable


- **Scale with table**

24-hour lifetime

Sub-second latency

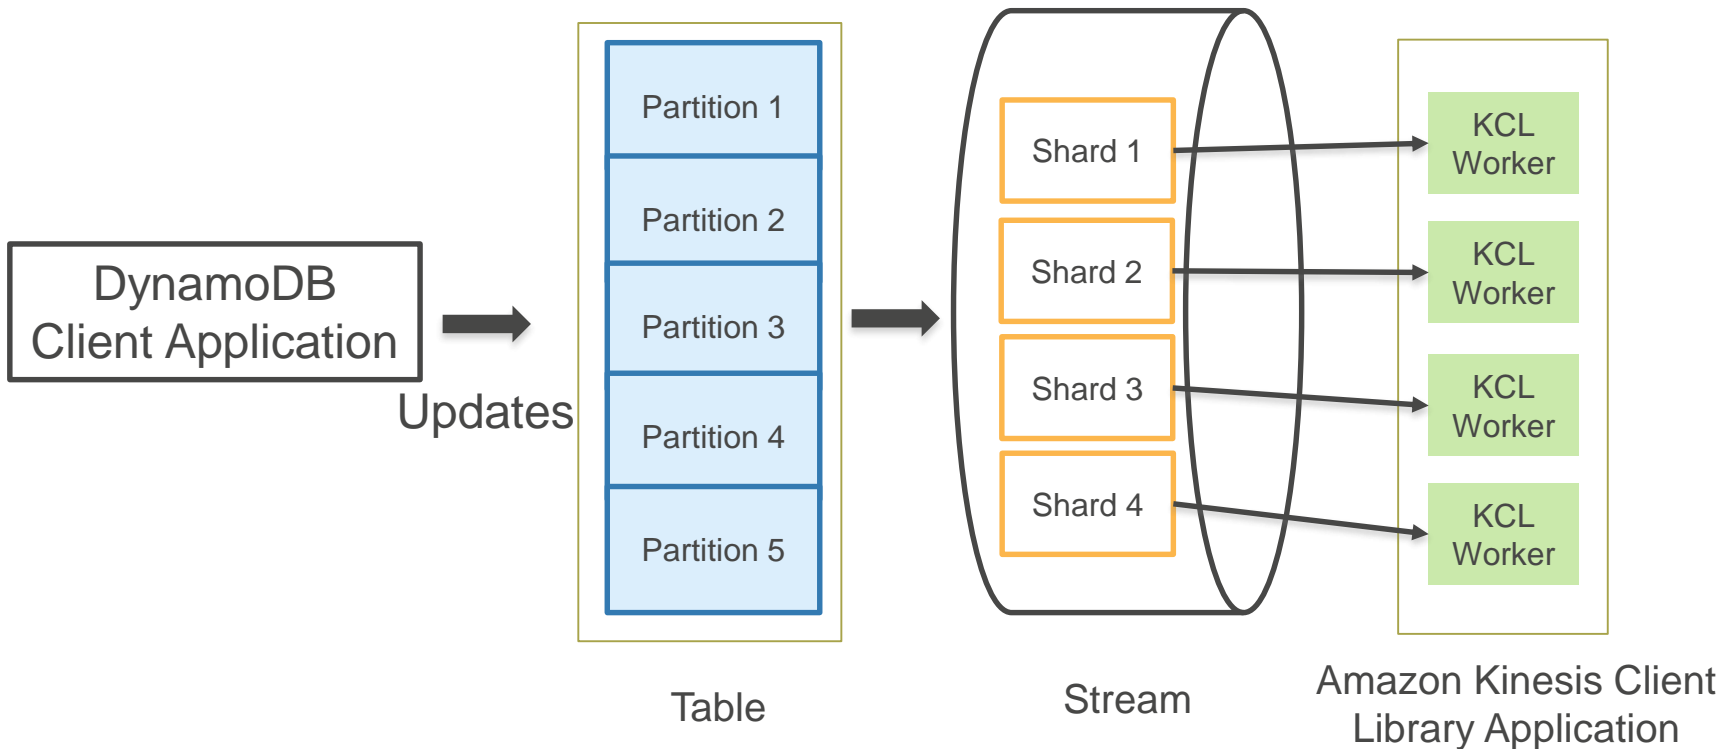
View types

UpdateItem (Name = John, Destination = Pluto)

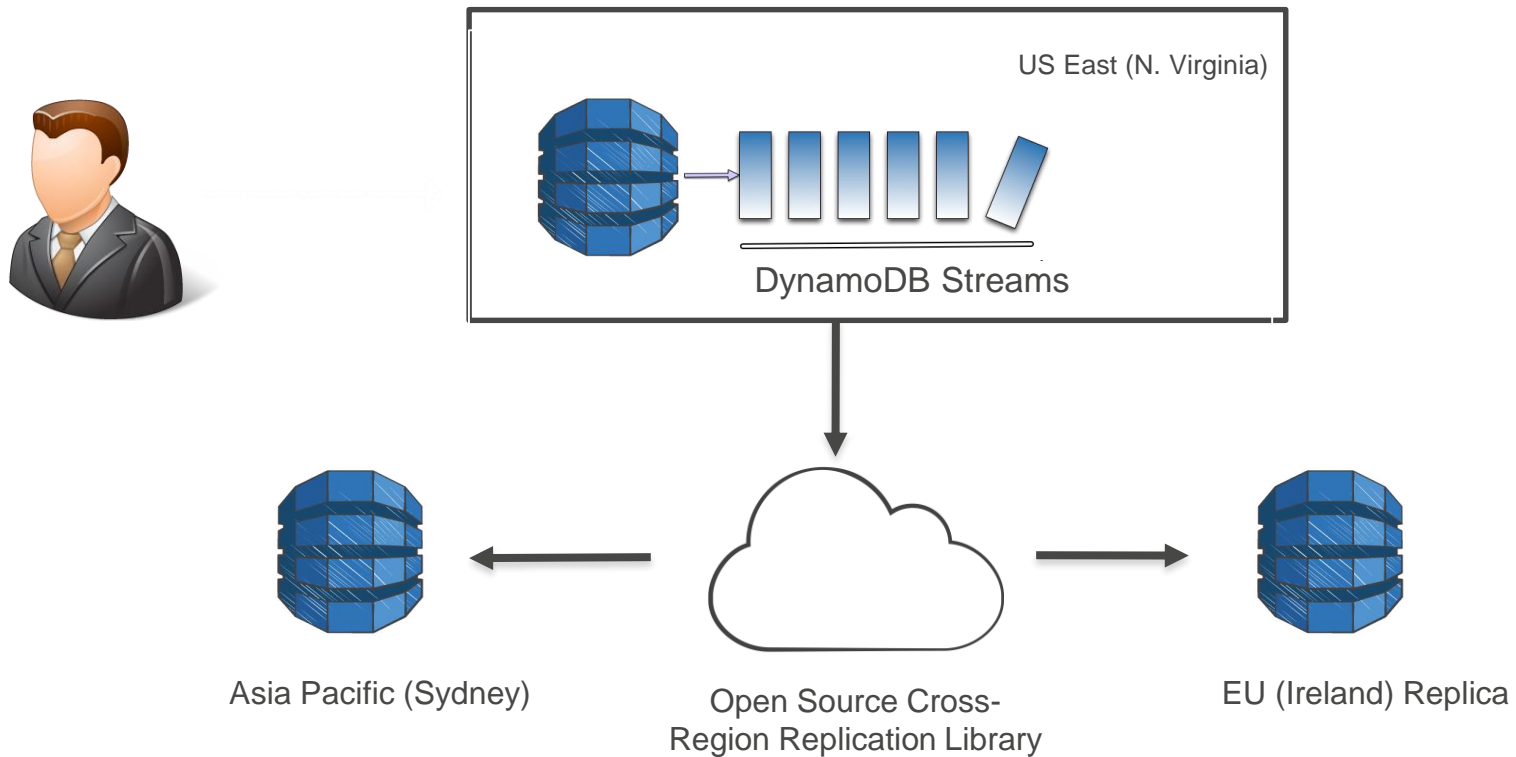


View Type	Destination
Old image—before update	Name = John, Destination = Mars
New image—after update	Name = John, Destination = Pluto
Old and new images	Name = John, Destination = Mars Name = John, Destination = Pluto
Keys only	Name = John

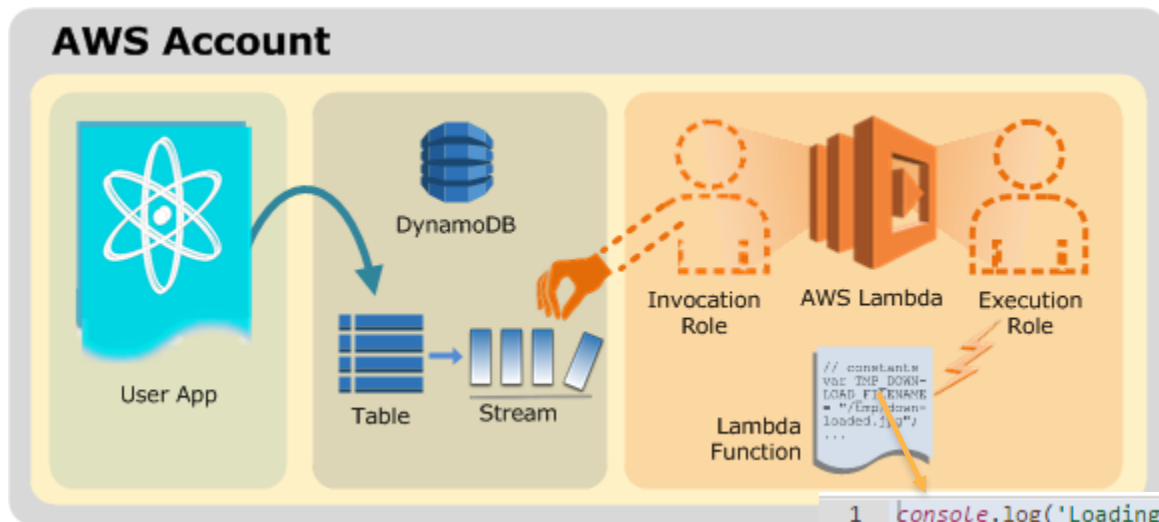
DynamoDB Streams and Amazon Kinesis Client Library



Cross-region replication



DynamoDB Streams and AWS Lambda



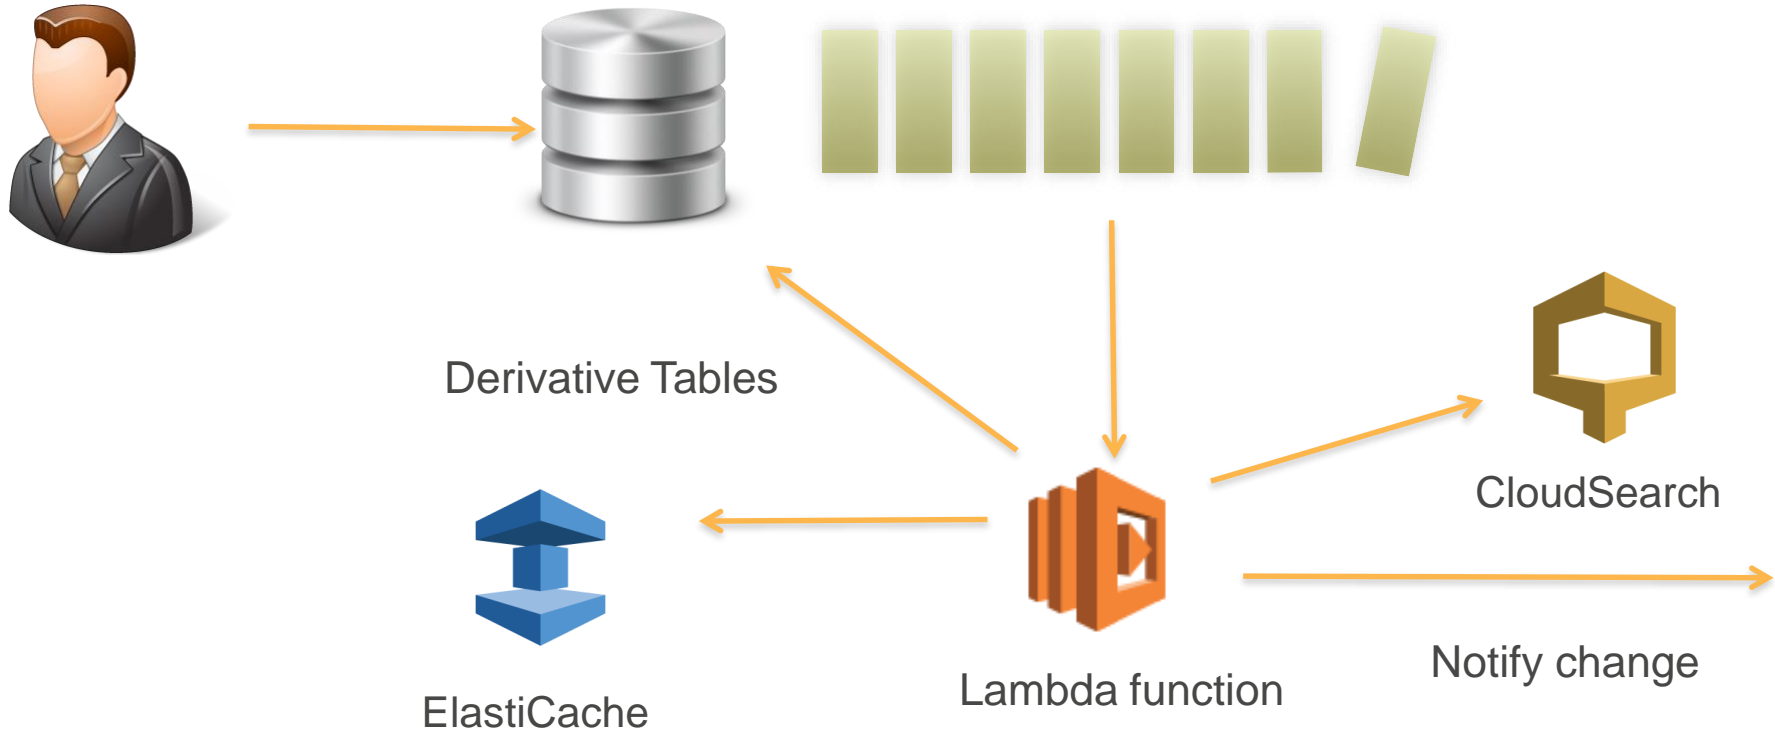
```
1 console.log('Loading event');
2 exports.handler = function(event, context) {
3   console.log("Event: %j", event);
4   for(i = 0; i < event.Records.length; ++i) {
5     record = event.Records[i];
6     console.log(record.EventID);
7     console.log(record.EventName);
8     console.log("DynamoDB Record: %j", record.Dynamodb);
9   }
10  context.done(null, "Hello World"); // SUCCESS with message
11 }
```

2015-03-21T07:44:58.883Z 2ca3769a-cf9e-11e4-b270-ad4d24b312ff INSERT

2015-03-21T07:44:58.883Z 2ca3769a-cf9e-11e4-b270-ad4d24b312ff DynamoDB Record: { "NewImage": { "name": { "S": "sivar" }, "hk": { "S": "3" } }, "SizeBytes": 15, "StreamViewState": "NEW_AND_OLD_IMAGES"


2015-03-21T07:44:58.883Z 2ca3769a-cf9e-11e4-b270-ad4d24b312ff Message: "Hello World"

Triggers






Analytics with DynamoDB Streams



Collect and de-dupe data in DynamoDB

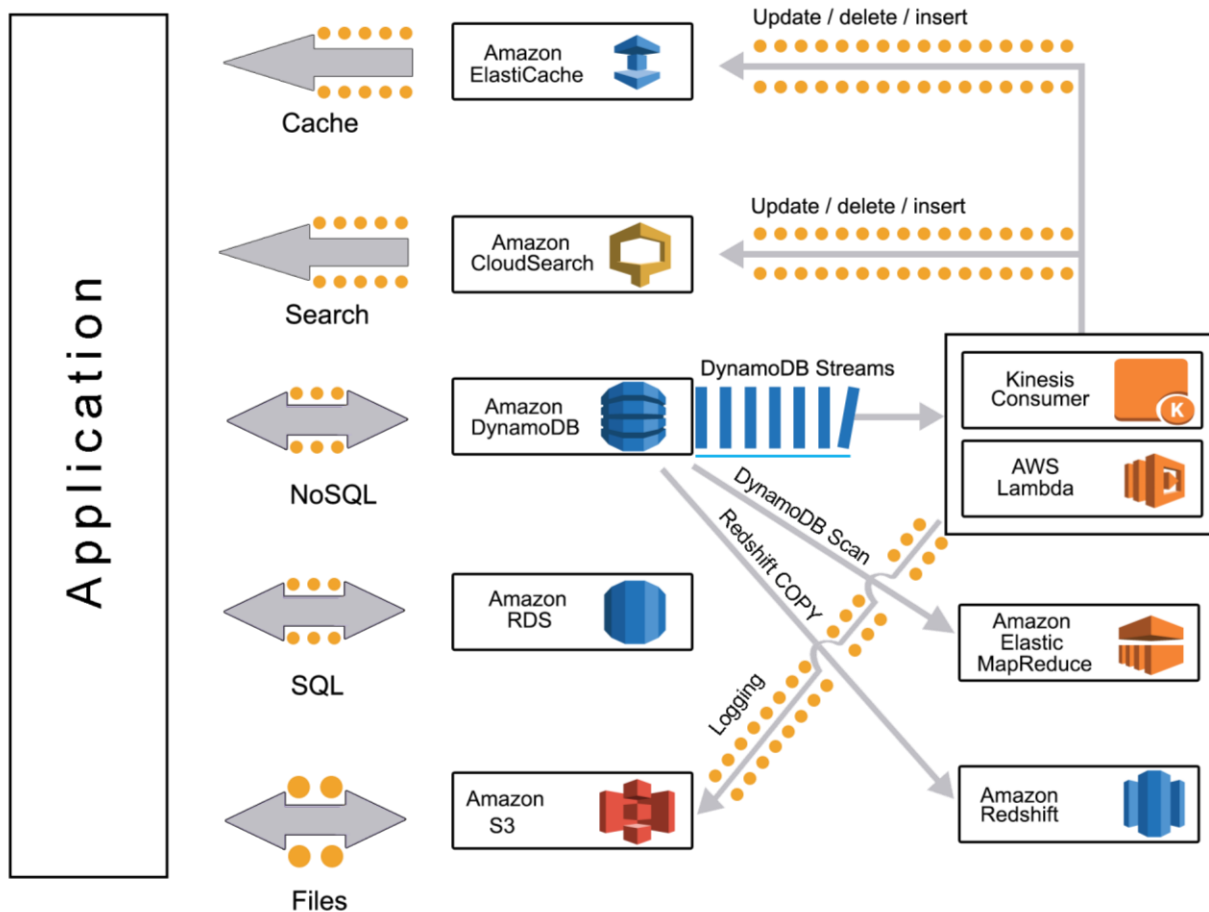
Aggregate data in-memory and flush periodically



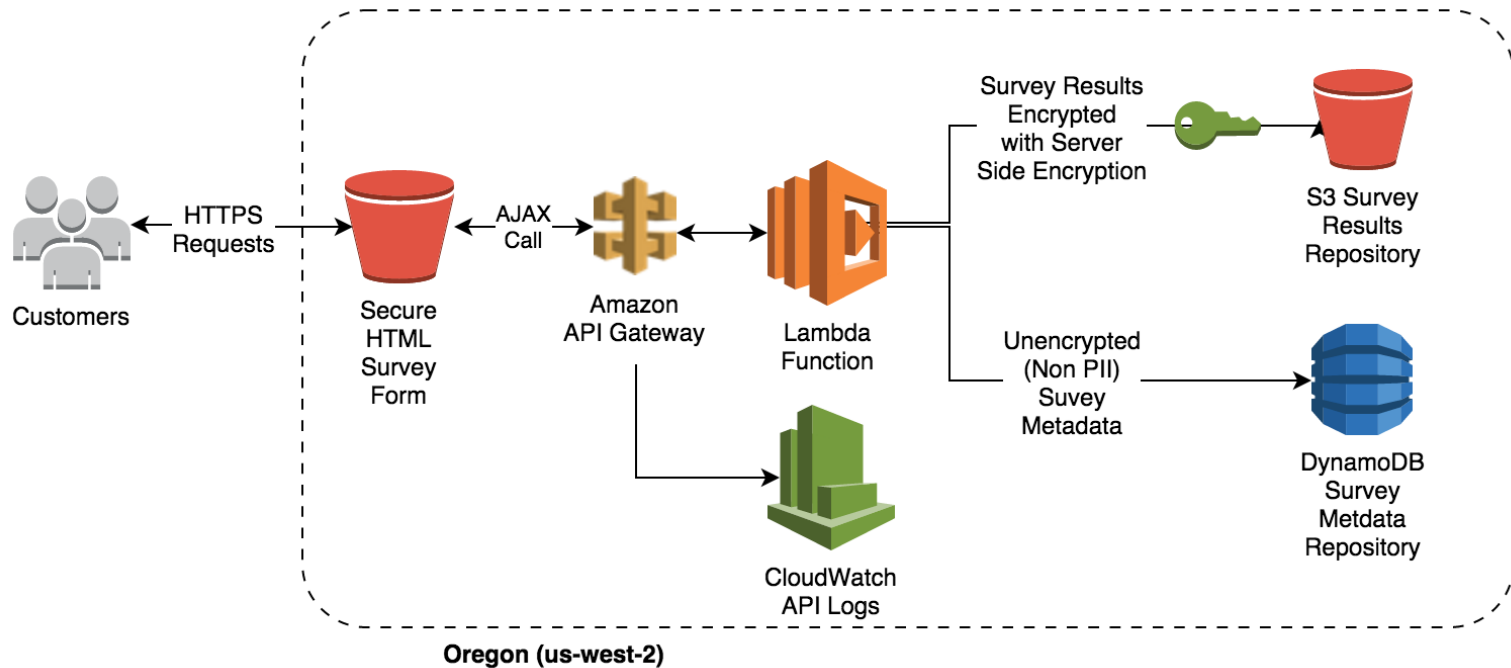
Important when: Performing real-time aggregation and analytics

Architecture

Reference Architecture



Elastic Event Driven Applications



Thank you!